Scalable Online-Offline Stream Clustering in Apache Spark

Anonymous
Anonymous Institution

Anonymous
Anonymous Institution

Abstract—Two of the most popular strategies to mine big data are distributed comput- ing and stream mining. The purpose of this thesis is to incorporate both together bringing a competitive stream clustering method into a modern framework for distributed computing, namely, Apache Spark. The method in question is CluS- tream, a stream clustering method which separates the clustering process into two different phases: an online phase which handles the incoming stream, generating statistical summaries of the data and an offline phase which takes those summaries to generate the final clusters. These summaries also contain valuable information which can be used for further analysis. The main goal is to adapt this method in such a framework in order to obtain a scalable stream clustering method which is open source and can be used by the Apache Spark community.

Keywords-big data streams; stream mining; stream clustering; CluStream

I. INTRODUCTION

The analysis of data streams comes along with important questions: what kind of data is it? What important information is contained in it? How does the stream evolve? The key question for this project among those is the latter, i.e. dealing with the evolution of the stream, because prior to the development of the CluStream [8] method there was not an easy to answer that question as it was one of the first to tackle this issue.

Clustering is one of the main tasks in data mining, also often referred as an exploratory subtask of it. As the name implies, the objective is to find clusters, i.e., collections of objects that share common properties. One can also relate this task to unsupervised machine learning, which intends to classify data when it lacks of labels, i.e., when the data instance does not indicate to which category it belongs. The CluStream method was developed in 2003 [8] and its main purpose is to pro- vide more information than previously developed algorithms for data stream clus- tering by that time. It provides a solution for handling streams of data indepen- dently from the one that finds the final clusters. It consists of two phases (passes) instead of one; the first one deals with the incoming data and stores relevant information over time and the second one is in charge of the clustering using the previously generated information.

II. RELATED WORK

CluStream has been implemented in different types of software and libraries, being *SAMOA - Scalable Advanced*

Massive Online Analysis one of the options. It is also a distributed computing implementation of the algorithm. The difference is that it is not implemented in *Spark*, but rather in a *Distributed Stream Processing Engine* which adapts the *MapReduce* approach to parallel stream processing [17].

Main differences with this adaptation:

- It does not include an offline macro-clustering phase.
- It is developed in *Java* and not designed to work with Spark.

StreamDM is a collection of algorithms for mining big data streams [1]. One of the included methods for stream clustering is *CluStream*. This collection of algorithms is developed for Spark.

Main differences with this adaptation:

- It does not include an offline macro-clustering phase.

III. BASIC CONCEPTS

A. CluStream

CluStream is a method developed in the Watson Research Center at IBM and the University of Illinois, UIUC. This method presented a different approach on the matter of clustering streams of data with respect to a modified version of *K-Means* which was adapted to work also with data streams. The main difference relies on the separation of the clustering process into two parts: one which would handle the data stream itself gathering only statistically relevant information (online part) and another which actually process the results of the former to produce the actual clusters wanted (offline part).

Separating the clustering process provides the user several advantages, among others:

- by saving only statistical data, rather than the original content, it is possible to save physical storage space (e.g. hard drive space) and therefore reducing costs and allowing a wider range in time to be clustered.
- The method also allows the analysis of the evolution of the data, as the necessary information for that is contained in the stored statistical information.
- Because the two parts operate independently it allows the user to select a time horizon, or even a time window, to perform the offline clustering part using the stored statistical information.

---

[1] As it is stated by them here: http://huawei-noah.github.io/streamDM/

1) The CluStream framework: This method is built over a few ideas that need to be conceptualized, which will answer fundamental questions and set up a basis of terminology useful along this work.

- **Micro-Clusters**: that is the given name for the statistical information summaries that is computed during the online component. They are a temporal extension of *cluster feature vectors* [18], which benefit from an additive feature that makes them a natural choice for the data stream problem [8].
- **Pyramidal time frame**: micro-clusters are stored periodically following a pyramidal pattern. This allows a nice tradeoff between the ability to store large amounts of information while giving the user the possible to work with different time horizons without loosing too much precision. The statistical summaries stored are used by the offline component to compute finally the macro-clusters which are the actual clusters the user intended to get.

2) Maintaining the micro-clusters: Whenever a new point arrives, it is necessary to find its nearest micro-cluster. It is possible to calculate an average radious or *RMSD*, only to then compare the distance to the point to a factor of it: when the distance between a point and its nearest micro-cluster is smaller or equal to the average radiuos (of the micro-cluster in question) times a user defined factor, then this point is added to the micro-cluster. Adding a point to a micro-cluster means that the properties of the micro-cluster change, such as RMSD and size (number of points).

Whenever a point (outlier) does not fulfill the mentioned condition, then a new micro-cluster has to be created in order to give this point a chance as a potential new cluster. In order to do so, an older micro-cluster has to be deleted or two micro-clusters have to be merged. To determine which solution is appropriate a recency value for each micro-cluster has to be determined[2] and until all the micro-clusters which have an older recency value than a user specified parameter are deleted, it is possible to start merging the micro-clusters which are closest to one another.

3) Offline macro-clusterig: The macro-clustering part is done by selecting a time window and then performing a modified version of *K-Means* to cluster the center of the current micro-clusters using the size as weights.

B. SPARK

Apache Spark is an open source framework developed in the AMPLab at the University of California. Traditionally, MapReduce and DAG engines are based on an acyclic data flow, which makes them non optimal for these applications listed above. In this flow, data has to be read from a stable storage system, like a distributed file system, and then processed on a series of jobs only to be written back to

---

[2]See [8] for more details.

---

the stable storage. This process of reading and writing data on each step of the workflow causes a significant rise in computational cost.

The solution proposed offers *resilient distributed datasets (RDDs)* to overcome this issue efficiently. RDDs are stored in memory between queries (no need of replication) and they can rebuild themselves in case of failure as they remember how they were originally built from other datasets by transformations such as *map, group, join*.

C. SPARK streaming

Figure 1: Flow of data in Spark streaming

Figure 1 shows the general idea of Spark streaming [4], a raw stream is linked to this module and it converts it to batches of data at user-defined intervals. These batches of data are then treated as RDDs, thus it gets distributed over the cluster where Spark runs. The abstraction of a data stream in Spark is called *DStream*, which stands for Discretized Stream, and is continuous series of RDDs. In a *DStream*, each RDD contains data from a specific interval of time, as it can be seen in Figure 2.

Figure 2: DStreams are Spark streaming's abstraction of a data stream

IV. CLUSTREAM IN APACHE SPARK

There are some modifications which had to be done in order to adapt *CluStream* in Spark. Working with Spark means worrking distributed computing and, thus, the algorithm has to be able to work in parallel. Both parts (online and offline) were adapted.

A. CluStreamOnline class (online phase)

Two processes were modified: processing the stream and updating the micro-clusters. As this adaptation uses Spark Streaming, the points coming from the stream are processed in batches at user specified time intervals. This contrasts with the original methodology which indicates to process point by point.

1) Finding nearest micro-cluster: The maintenance of the micro-clusters starts with this operation. After initialization (described in [8]) is performed, finding the nearest micro-clusters for all the points is the very first thing to be done for every new batch of data.

Algorithm 1 Find nearest micro-cluster.

Input: *rdd: RDD[breeze.linalg.Vector[Double]], mcInfo: Array[(McI,Int)]—* *rdd* is an RDD containing data points and *mcInfo* is the collection of the micro-clusters information.

Output: *rdd: RDD[(Int, breeze.linalg.Vector[Double])]* — returns a tuple of the point itself and the unique ID of the nearest micro-cluster.

1: **for all** $p \in rdd$ **do**
2:     $minDistance \leftarrow Double.PositiveInfinity$
3:     $minIndex \leftarrow Int.MaxValue$
4:     **for all** $mc \in mcInfo$ **do**
5:         $distance \leftarrow squaredDistance(p, mc_1.centroid)$
6:         **if** $distance \leq minDistance$ **then**
7:             $minDistance \leftarrow distance$
8:             $minIndex \leftarrow mc_2$
9:         **end if**
10:     **end for**
11:     $p = (minIndex, p)$
12: **end for** **return** rdd

Finding the nearest micro-clusters is an operation of complexity $O(n * q * d)$, where n is the number of points, q the number of micro-clusters and d the dimension of the points; q and d remain constant during runtime but n might vary. Algorithm 1 describes a simple search for the minimum distance for every point in the RDD to the micro-clusters. This is also a good opportunity to show how this works using Spark and *Scala*:

```
1  def assignToMicroCluster(rdd: RDD[Vector[Double]],
       mcInfo: Array[(MicroClusterInfo, Int)]): RDD
       [(Int, Vector[Double])] = {
2    rdd.map { a =>
3      var minDist = Double.PositiveInfinity
4      var minIndex = Int.MaxValue
5      for (mc <- mcInfo) {
6        val dist = squaredDistance(a, mc._1.
             centroid)
7        if (dist < minDist) {
8          minDist = dist
9          minIndex = mc._2
10       }}
11     (minIndex, a)
12   }}
```

Spark uses this *map* operation to serialize the function passed so that all nodes in the cluster get the same instruction, this is exactly how computations are parallelized within this framework. At this point, every node performs this operation to find the nearest micro-cluster for all the points they locally have.

2) Processing points: The points are separated in two: points within micro-clusters and outliers, it is possible to compute the necessary information from them to update the micro-clusters. It is important to perform this step before handling the outliers because this adaptation process the points for batches of data and not points individually as they arrive, and the reasons are:

- Every point in a batch is treated equally in terms of temporal properties as this batch gets distributed among the cluster with the same time stamp and there is no constant communication among nodes.
- The process of handling outliers involves deleting and merging micro-clusters for every outlier, modifying the micro-clusters' structure. A sequential implementation was preferred to reduce communication costs.

These two points are some of the key differences between the original $CluStream$ method and this adaptation. For the original, it is possible to handle point by point as each have different clear time stamps.

3) Handling outliers: First the micro-clusters which are safe to delete are determined, then the outliers can be handled. In general, there are three possible scenarios for outliers:

- If the point lies within the restriction regarding the RMSD for its nearest micro-cluster in the array $newMicroClusters$, the point is added to it. This stores all newly created micro-clusters.
- If the point does not lie within any of the new micro-clusters, then it replaces a micro-cluster from the $safeDelete$ array, while there are safe-to-delete micro-clusters.
- If none of the previous scenarios are viable, then the two micro-clusters that are closest to each other get merged, freeing one spot to create the new micro-cluster. This is the the most computationally expensive scenario. The function $getTwoClosestMicroClusters()$ has a complexity of $O(p_m d \cdot \frac{q!}{2!(q-2)!})$, where p_m is the number of outliers that require a merge, d the dimension of the points, and q the number of micro-clusters.

It is important in the procedure described in Algorithm 2 to locally update the $mcInfo$ every time a point is added to a micro-cluster, two micro clusters are merged and when a new micro-cluster is created. There could be a lot of change, depending on the outliers, and this loop requires up-to-date information for each iteration, otherwise merges and the RMSD check would be inaccurate.

B. CluStream class (offline phase)

Using a *weighted K-Means* approach, as described in [8] was not directly possible, and for that reason, a new adaptation had to be done in order to achive similar results.

Algorithm 2 handle outliers.

1: $j \leftarrow 0$
2: **for all** $p \in dataOut$ **do**
3:     $distance, mcID \leftarrow$
    $getMinDistanceFromIDs(newMicroClusters, p_2)$
4:     **if** $distance < t * mcInfo[mcID]_1.rmsd$ **then**
5:         $addPointToMicroCluster(mcID, p_2)$
6:     **else**
7:         **if** $safeDelete[j].isDefined$ **then**
8:             $replaceMicroCluster(safeDelete[j], p_2)$
9:             $newMicroClusters.append(j)$
10:             $j \leftarrow j + 1$
11:         **else**
12:             $index1, index2 \leftarrow$
    $getTwoClosestMicroClusters(keepOrMerge)$
13:             $mergeMicroClusters(index1, index2)$
14:             $replaceMicroClusters(index2, p_2)$
15:             $newMicroClusters.append(j)$
16:             $j \leftarrow j + 1$
17:         **end if**
18:     **end if**
19: **end for**

The fakeKMeans solution:

The original *CluStream* method suggests to use a slightly modified version of K-Means, a version for which one can initialize the seeds (initial clusters) by sampling from the micro-clusters' centroids taking into account the number of points each micro-cluster has and for which one can use these centroids as weighted input points. These weights, again, are related to the number of points they absorbed. Spark's (current) implementation of K-Means does allow to initialize the seeds but unfortunately it is not possible to predefine the weights for the input points.

In order to solve this issue, a new version of K-Means needs to be implemented. This version uses, in fact, Spark's own version, but to overcome the problem of not being able to define the weights at the beginning, this new version uses as input many points sampled from the micro-clusters' centroids.

V. Experiments

A. Experiments setting

For the experiments we used the Network Intrusion dataset[3], which consists of 494,021 instances. For the analysis, we used only the numerical attributes (#34 out of #43 attributes). We vary the speed of the stream and the horizon and we derive two different stream configurations. The first one, denoted as $DS1$, has a speed $v = 2,0000$ points per timestamp, which implies that it lasts for $\frac{494,021}{2,000} \approx 247$

---

[3]Source: http://kdd.ics.uci.edu/databases/kddcup99/kddcup99.html

---

time units, and a horizon $H = 1$. The second one, denoted as $DS2$, has a speed of $v = 200$ points per timestamp, thus lasting for a period of $\frac{494021}{200} \approx 2470$ time points, and a horizon $H = 256$. To evaluate the clustering quality, we report in Section V-B on the sum of square distances (SSQ) from the points to their nearest micro-cluster, using Euclidean distance as the distance function, within a horizon H. Efficiency and scalability is discussed in Section V-C.

B. Clustering quality

We first compare *CluStream-Spark* to the original *CluStream* (Section V-B1) and then against other stream clustering approaches in Spark (Section V-B2).

1) Quality of $Spark - CluStream$ vs original CluStream:

Results for DS1: The SSQ for $DS1$ is shown in Figure 9 for the original $CluStream$ and our $Spark - CluStream$. We used the same parameters as in [8], i.e., $\alpha = 2, l = 10, InitNumber = 2000, \delta = 512, t = 2$. The parameter m, for m last points, was the only one not provided, we set it to $m = 20$. m is used to determine the approximate recency value as if the time of arrival of the last m points was averaged. For $DS1$, both m and δ are irrelevant and the reason is that the threshold is never reached (247 time units vs. 512). The number of micro-clusters was set to $q = 50$, 10 times the number of final clusters (5).

In Table I we show the SSQ scores for our *CluStream-Spark* and the results from *CluStream* as reported in [8] - these are approximate results "extracted" from the original paper. Although we don't know the exact values for *CluStream*, we can see that the magnitudes of the average SSQ are comparable.

Results for DS2: Parameters were set as for $DS1$. The exact scores of *CluStream-Spark* and the approximate values from *CluStream* are reported in Table II. Again the quality scores are comparable, i.e., our *CluStream-Spark* achieves similar quality clusterings to the original *CluStream* .

2) $Spark - CluStream$ vs other clustering approaches in SPARK: We compare our *CluStream-Spark* against available solutions for stream clustering in Spark and in particular against *Streaming K-Means* [4] and *StreamDM-CluStream* [5]. We report here on their clustering quality, the efficiency issue is discussed in Section V-C.

We roughly overview these methods hereafter.

Results on $DS1$: For the $DS1$ stream the results are shown in Figure 3. For *Streaming K-Means*, the horizon $H = 1$ is transformed to $halfLife = 1,000$ points, as the stream speed is 2,000 points. The *decayFactor* is set to 0, i.e., only the last 2,000 points will influence on the clusters. *StreamDM-CluStream* is set up with its default parameters,

---

[4]More information: https://databricks.com/blog/2015/01/28/introducing-streaming-k-means-in-spark-1-2.html

[5]More information: http://huawei-noah.github.io/streamDM/

$DS1$ - avg SSQ	10k	40k	160k	320k
$CluStream$	10^5-10^6	10^{12}-10^{13}	$\approx 10^6$	10^2-10^3
$Spark - CluStream$	3.099×10^5	6.676×10^{12}	7.833×10^5	4.191×10^2

Table I: $DS1$ - Average SSQ values

$DS2$ - avg SSQ	150k	250k	350k	450k
CluStream	10^{13}-10^{14}	$\approx 10^5$	10^{12}-10^{13}	$\approx 10^8$
Spark-CluStream	5.402×10^{13}	5.143×10^4	1.892×10^{13}	9.646×10^7

Table II: $DS2$ - Average SSQ values

only changing the horizon to 1 and the number of micro-clusters to 50 to match those of *CluStream-Spark* . As we can

Figure 3: Average SSQ for different stream clustering methods in SPARK. $DS1$ (Stream speed $v = 2,000$, H=1)

see, our *CluStream-Spark* delivers results which are very close to those of *Streaming K-Means*, whereas *Streaming K-Means* with the *decayFactor* (DF) is the best. The surprising results come from *StreamDM-CluStream*, which is the worse among the tested methods, especially for the last two points, i.e., at $160k$ and $320k$. To further investigate this, we conduct another experiment where we compare *CluStream-Spark* without snapshots against *StreamDM-CluStream*. The results are shown in Figure 4.

As we can see, *CluStream-Spark* outperforms *StreamDM-CluStream* even if we remove the snapshot part, but the quality is lower comparing to the snapshot version.

Results on $DS2$: For the $DS2$ stream the results are shown in Figure 5. All parameters remained the same for all methods, except for the $halfLife$ parameter for *Streaming K-Means*, which is set to $halfLife = 25,600$ ($200 \cdot 256 = 51,200$). We calculate the $decayFactor$ as follows: At 150,000 points, the ratio of the points to cluster to the total number of points at that particular time is $\frac{51,200}{150,000} \approx 0.3413$. At 250,000 points, this equals to $\frac{51,200}{250,000} \approx 0.2048$. At 350,000 points, this equals to $\frac{51,200}{350,000} \approx 0.1462$. At 450,000 points, this equals to $\frac{51,200}{450,000} \approx 0.1137$. Averaging


Figure 4: *CluStream-Spark* without snapshots vs *StreamDM-CluStream*. $DS1$ (Stream speed $v = 2,000$, $H = 1$, $m = 100$).

those ratios leads to a *decayFactor = 0.2015*. As we can

Figure 5: Average SSQ for different stream clustering methods in SPARK. $DS2$ (Stream speed $v = 200$, H=256)

see, *CluStream-Spark* performs consistently well and better than *StreamDM-CluStream*. As expected *Streaming K-Means* with the *decayFactor* achieves the best performance.

We repeat the without-snapshot experiment, the results are shown in Figure 6. As we can see, *CluStream-Spark* delivers better results than *StreamDM-CluStream* but the difference is reduced significantly comparing to $DS1$. These results might indicate that *StreamDM-CluStream* benefits from larger horizons.
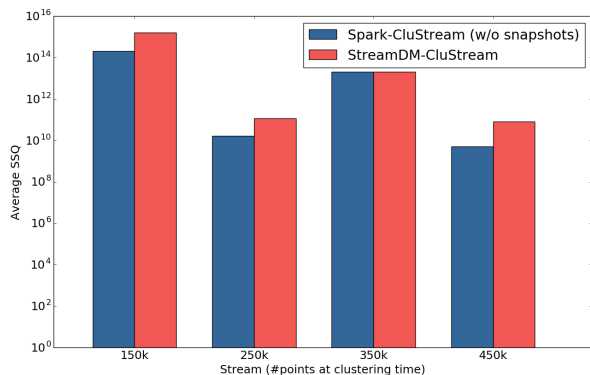
Figure 6: *CluStream-Spark* without snapshots vs *StreamDM-CluStream*. $DS2$ (Stream speed $v = 200$, $H = 256$, $m = 100$).

C. Scalability

We test the scalability with respect to data dimensionality and number of microclusters, using data generated by a *Random Radial Basis Function* generator.

The scalability tests are performed in two different scenarios: one being an analysis of how it scales for different number of attributes (dimensions of the data points) using only 20 micro-clusters and the other one using 200 micro-clusters as the number of attributes and the number of final clusters for a specific purpose are two key factors which determine the complexity of *Spark-CluStream*. The speed of the stream is controlled for 10000 points for every batch of data because it is easier to test the scalability when many computations have to be done. An application using Spark streaming assigns one core exclusively to handle the stream, therefore the number of processors mentioned in here is the total, but the real number of processors used for the computations is that number minus one.

Figure 7 shows that using only 20 micro-clusters and 2 dimensions has poor scalability, not even being able to perform twice as fast as for a single processor (2 in total). Even for this high speed streaming, one processor is enough to process the batches of data before a new batch is processed, meaning that the setup is stable.

Increasing the dimensionality of the points increases the computational effort needed to process the points in every batch of data and here is where *Spark-CluStream* shows its scalability, which is almost linear[6] for up to 16-17 processors, as it can be seen in Figure 8. From the average processing time per batch, it can be seen that from 32 to 40 processors it does not improve much anymore and the speedup does not increase quasi-linearly anymore. Here

---

[6]By linear scalability does not mean it scales with a 1 to 1 ratio, but rather linearly proportional.

a total of 9 processors were required to stabilize *Spark-CluStream*.

Interestingly, increasing the number of micro-clusters by a factor of 10 for 2 attributes resulted in good scalability, similarly to the scenario with 20 micro-clusters and 100 attributes. Here a total of 8 processors were enough for a stable run, as shown in Figure 9.

Finally, when the number of clusters and the number of attributes are both increased significantly, Figure 10 shows for *Spark-CluStream* quasi-linear scalability but this time only up to about 8-9 processors. After that point, the speedup slows down showing almost no improvement after 16 processors. This test never reached a stable configuration.

D. Performance

In this section, the scalability of *Spark-CluStream* is compared to that of *StreamDM-CluStream* and Spark's *Streaming K-Means* using the Spark cluster setup for $q = 20$ and $d = 100$, for the $CluStream$ method. Also, a test on a single machine is performed.

When it comes to higher dimensions, *Spark-CluStream* shows a significant improvement over *StreamDM-CluStream*, which never got to the point were it was stable (below the 1 second mark), as shown in 11, it seems to scale as fast as *Spark-CluStream* but it was not enough even with 40 processors.

Another interesting comparison, is the processing time per batch of data for a single machine, using a real dataset such as the *Network Intrusion*. Here, communication is less of an issue as all the partitions lie in the same share memory space, and still there are 4 virtual cores in disposition for the algorithms to run.

The test was performed using a stream speed of 2000 points per batch and with a horizon $H = 1$, to match one of the validation tests.

The results shown in Figure 12 are quite remarkable. As *StreamDM-CluStream* shows a very significant disadvantage when using greater numbers of micro-clusters and higher dimensions.

For this single machine test, *Spark-CluStream* was about 18 times faster on average than *StreamDM-CluStream* and about two times slower than *Streaming K-Means* on average.

Another consideration to be made, is that *Spark-CluStream* saves a snapshot for every batch of data, having to write to disk, while the other two algorithms never access the disk for this matter.

VI. CONCLUSIONS

We proposed a variation of *CluStream* tailored to Apache Spark, *CluStream-Spark* . Our experiments show that *CluStream-Spark* achieves similar quality to the original approach [8], while it is far more efficient. Comparing to other stream clustering approaches in Spark, our results show that *Streaming K-Means* is the fastest algorithm among

Figure 7: Dimension: $d = 2$

Figure 8: Dimension: $d = 100$

Scalability-dimensionality comparison for Stream speed = 10000 and q = 20


Figure 9: Dimension: $d = 2$

Figure 10: Dimension: $d = 100$

Scalability-dimensionality comparison for Stream speed = 10000 and q = 200

the three tested (highly optimized for Spark), but it does not offer the flexibility of the online-offline clustering approaches like *CluStream* that better fit evolving data streams. Comparing *CluStream-Spark* to *StreamDM-CluStream*, *CluStream-Spark* delivers more consistent and accurate results, and outperforms *StreamDM-CluStream* in most cases, including one up to around 18 times faster.

As part of our future work, we will focus on minimizing the communication cost between the different nodes to further improve the efficiency and also on dealing with outliers to improve quality.

REFERENCES

[1] Apache spark. http://spark.apache.org/. Accessed: 08.07.2015.

[2] Apache spark: Preparing for the next wave of reactive big data. https://www.typesafe.com/blog/apache-spark-preparing-for-the-next-wave-of-reactive-big-data. January 27, 2015.

[3] Apache spark, research. http://spark.apache.org/research. Accessed: 24.01.2016.

[4] Apache spark, streaming. http://spark.apache.org/docs/latest/streaming-programming-guide.html. Accessed: 01.03.2016.

[5] Databricks, streaming k-means. https://databricks.com/blog/2015/01/28/introducing-streaming-k-means-in-spark-1-2.html. Accessed: 01.03.2016.

[6] Scikit-learn: Mini batch kmeans. http://scikit-learn.org/stable/modules/clustering.html. 2010-2014 BSD License, Accessed: 08.07.2015.

[7] Spark's mlib: Clustering. http://spark.apache.org/docs/latest/mllib-clustering.html. Accessed: 08.07.2015.

[8] Charu C. Aggarwal, Jiawei Han, Jianyong Wang, and Philip S. Yu. A framework for clustering evolving data streams. In *Proceedings of the 29th International Conference on Very Large Data Bases - Volume 29*, VLDB '03, pages 81–92. VLDB Endowment, 2003.

Figure 11: Processing time comparison: $q = 20$, $d = 100$

Figure 12: Processing time comparison for a single machine: $q = 50$, $d = 34$

ACM Symposium on Operating Systems Principles, SOSP '03, pages 29–43, New York, NY, USA, 2003. ACM.

[14] Satish Gopalani and Rohan Arora. Article: Comparing apache spark and map reduce with performance analysis using k-means. *International Journal of Computer Applications*, 113(1):8–11, March 2015. Full text available.

[15] Hai-Guang Li, Gong-Qing Wu, Xue-Gang Hu, Jing Zhang, Lian Li, and Xindong Wu. K-means clustering with bagging and mapreduce. In *System Sciences (HICSS), 2011 44th Hawaii International Conference on*, pages 1–8, Jan 2011.

[16] SimoneA. Ludwig. Mapreduce-based fuzzy c-means clustering algorithm: implementation and scalability. *International Journal of Machine Learning and Cybernetics*, pages 1–12, 2015.

[17] Gianmarco De Francisci Morales and Albert Bifet. Samoa: Scalable advanced massive online analysis. *Journal of Machine Learning Research*, 16:149–153, 2015.

[18] Tian Zhang, Raghu Ramakrishnan, and Miron Livny. BIRCH: an efficient data clustering method for very large databases. In Jennifer Widom, editor, *SIGMOD '96: Proceedings of the 1996 ACM SIGMOD International Conference on Management of Data*, pages 103–114. ACM Press, 1996.

[9] E. Alpaydin. *Introduction to Machine Learning*. Adaptive Computation and Machine Learning. MIT Press, 2014.

[10] Christopher M. Bishop. *Pattern Recognition and Machine Learning (Information Science and Statistics)*. Springer-Verlag New York, Inc., Secaucus, NJ, USA, 2006.

[11] Jeffrey Dean and Sanjay Ghemawat. Mapreduce: simplified data processing on large clusters. In *OSDI 04: PROCEEDINGS OF THE 6TH CONFERENCE ON SYMPOSIUM ON OPERATING SYSTEMS DESIGN AND IMPLEMENTATION*. USENIX Association, 2004.

[12] D. Garg and K. Trivedi. Fuzzy k-mean clustering in mapreduce on cloud based hadoop. In *Advanced Communication Control and Computing Technologies (ICACCCT), 2014 International Conference on*, pages 1607–1610, May 2014.

[13] Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung. The google file system. In *Proceedings of the Nineteenth*